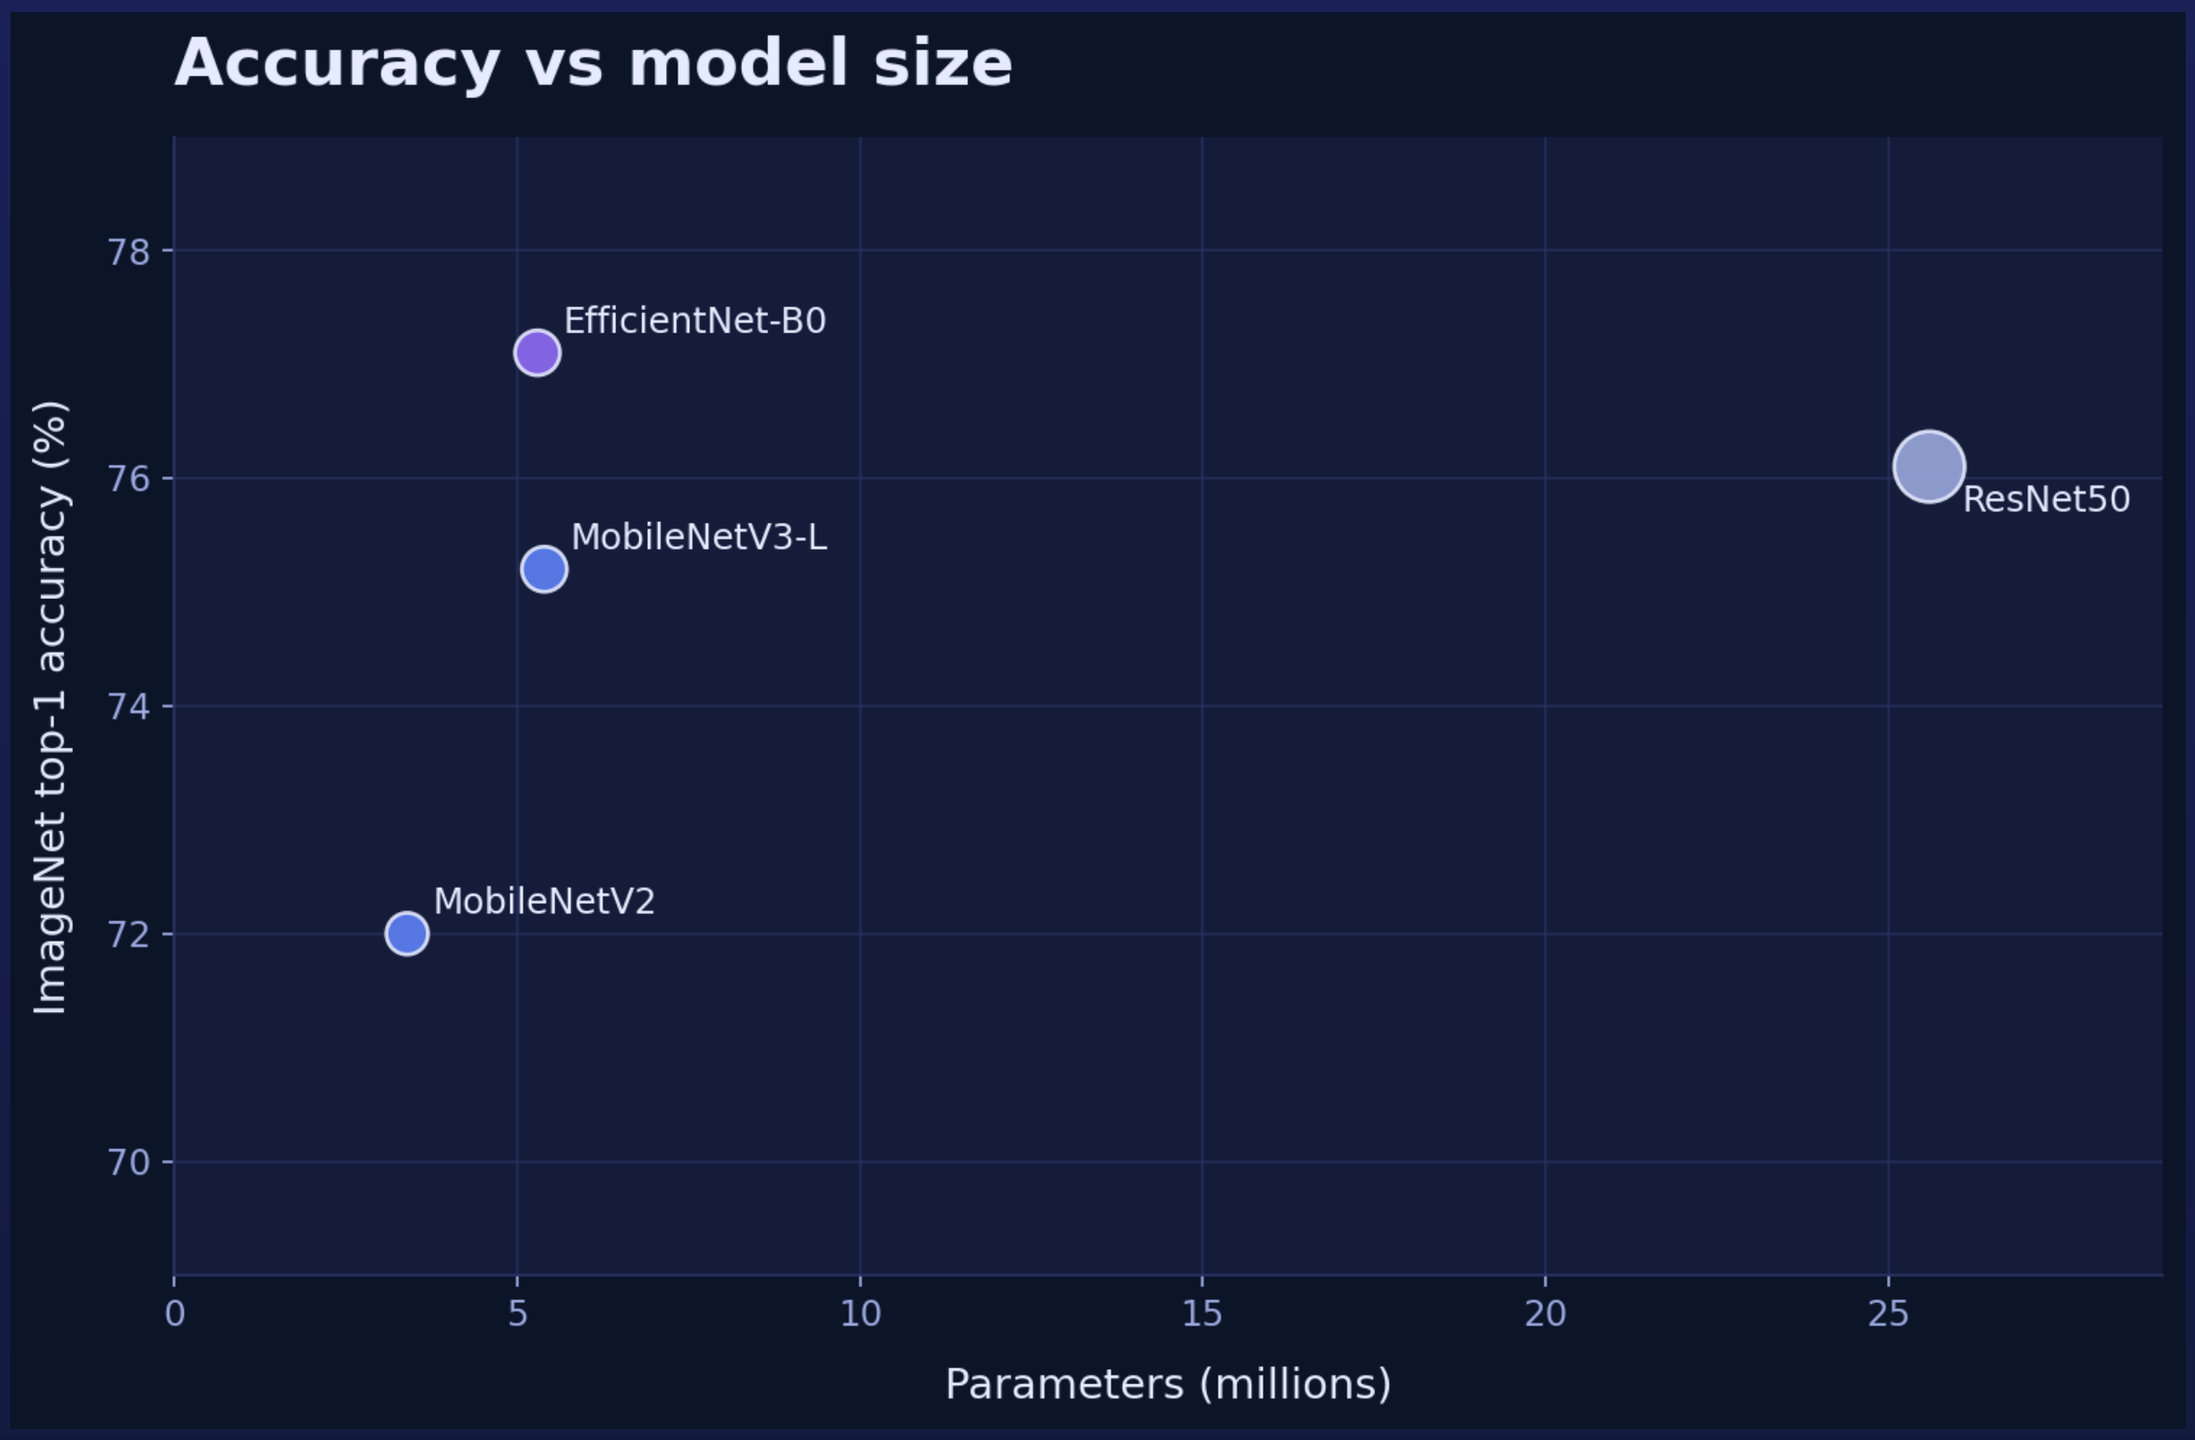
MobileNet Makes Computer Vision More Practical on Limited Hardware

ARTIFICIAL INTELLIGENCE · MACHINE LEARNING · DEEP LEARNING

MobileNet Makes Computer Vision More Practical on Limited Hardware

SWIPE ->

THE NUMBERS



MobileNetV2 reaches 72% with 3.4M parameters - a fraction of ResNet50's 25.6M.

Not every computer vision application runs on a powerful GPU.

Some need to run on mobile phones, edge devices, embedded systems, or hardware with strict power constraints.

That's where MobileNet becomes interesting.

One of its key ideas is the depthwise separable convolution.

A standard convolution combines spatial filtering and channel mixing in one operation.

A depthwise separable convolution splits that work into two stages:

- Depthwise convolution: Apply spatial filters separately to each input channel.
- Pointwise convolution: Use 1x1 convolutions to combine information across channels.

This can reduce computation and parameter requirements substantially compared with a standard convolution.

Why does that matter?

Lower computational cost can make inference more practical on resource-constrained devices.

However, fewer operations don't guarantee lower real-world latency on every device.

Hardware acceleration, memory movement, runtime implementation, and input resolution all matter.

When choosing a vision model for deployment, I would benchmark the actual target device instead of relying only on theoretical FLOPs.

A model is only efficient if it performs efficiently where you need to run it.

#MobileNet #EdgeAI #ComputerVision

Thanks for reading.

Follow **Yasir Usman** for more on AI, ML and computer vision.

Code: github.com/yasirusman85/Tank