Modern OCR Is Much More Than Reading Text

ARTIFICIAL INTELLIGENCE · MACHINE LEARNING · DEEP LEARNING

Modern OCR Is Much More Than Reading Text

SWIPE ->

OCR used to be explained as a simple task:

Image in text out.

That explanation is no longer enough for many real-world applications.

Consider a scanned invoice.

Knowing that the document contains the words "Total Amount" and "$2,500" isn't sufficient.

The system must also understand that the amount belongs to the total field, not the invoice number or an unrelated line item.

A modern document-processing pipeline may involve:

- Image preprocessing and deskewing
- Text detection and recognition
- Document layout analysis
- Table and key-value extraction
- Reading-order reconstruction
- Validation and structured output

Depending on the use case, this can involve traditional OCR engines, deep learning recognition models, vision-language models, or a combination of them.

But here's the important part:

Better text recognition doesn't automatically mean better document understanding.

A system can read every word correctly and still assign the wrong value to the wrong field.

For production OCR, I would evaluate both character-level accuracy and whether the final extracted data is actually correct.

Because the business doesn't need a document full of recognized words.

It needs reliable information.

#OCR #DocumentAI #ComputerVision

Thanks for reading.

Follow **Yasir Usman** for more on AI, ML and computer vision.

Code: github.com/yasirusman85/Tank